

I.4: Numpy

- One of the most important foundational packages for numerical computing in Python
- It provides efficient multidimensional array type `ndarray` with fast array-oriented arithmetic operations.
- It allows executing fast operations on entire arrays of data without having to write loops.
- It provides Linear algebra operations, random number generation, and Fourier transform capabilities.
- it is designed for efficiency on large arrays of data.
- NumPy internally stores data in a contiguous block of memory, independent of other built-in Python object
- It is a computational foundation for general numerical data processing for many other packages such as `pandas` and `SciPy`.

```
import sys
import numpy as np

py_arr = list(range(1000000)) #standard python array
np_arr = np.arange(1000000) #numpy array

#print time consumed in multiplicatios
%time for _ in range(10): np_arr * 2/2
%time for _ in range(10): [x * 2/2 for x in py_arr]

CPU times: user 31.8 ms, sys: 6.08 ms, total: 37.9 ms
Wall time: 51.8 ms
CPU times: user 1.57 s, sys: 180 ms, total: 1.75 s
Wall time: 2.26 s
```

NumPy-based algorithms are generally 10 to 100 times faster (or more) than their pure Python counterparts and use significantly less memory.

NumPy ndarray (Multidimensional Array)

- a key feature of NumPy is its N-dimensional array object, or `ndarray`.
- It is a fast, flexible container for large datasets in Python.
- It enables you to perform mathematical operations on whole blocks of data similarly to the operations between scalars.
- It is a generic multidimensional container for homogeneous data(same type).
- It has `shape`, a tuple indicating the size of each dimension, and `dtype`, an object describing the data type of the array

```
import numpy as np
np_data = np.random.randn(4,3) # create an array of random values with
dims 4x3
print("np_data: \n", np_data)

#do some math operations with the array
np_data = (np_data + np_data) * 2
```

```

print("\n(np_data + np_data) * 2: \n", np_data)

#print shape and dtype of the array
print("\n shape: ", np_data.shape, "; dtype: ", np_arr.dtype)

np_data:
[[-0.97964204 -0.25464528  1.14784297]
 [-0.19705887 -1.20501532  0.6985225 ]
 [-0.29363815 -0.35362777 -0.67977467]
 [-1.11033485  1.26349186  0.63231134]]

(np_data + np_data) * 2:
[[-3.91856818 -1.01858113  4.59137189]
 [-0.78823548 -4.82006128  2.79409   ]
 [-1.17455261 -1.41451109 -2.71909869]
 [-4.44133941  5.05396746  2.52924536]]

shape: (4, 3) ; dtype: int64

```

Basic operations on ndarray

The easiest way to create an array is to use `array` function

```

arr1 = np.array([1.5, 5.1, 6, 12, 3.4])
#create an array 1d
print('\n arr1:\n', arr1)

arr2 = np.array([arr1, arr1, [0, 1, 2, 3, 8]])
#create an array 2d
print('\n arr2:\n', arr2)

arr3 = np.array([arr1, [0.1, 1.2, 2.3, 3.4, 8.9]], dtype= np.int32)
#create an array given the type
print('\n arr3:\n', arr3)

arr1:
[ 1.5  5.1  6.  12.  3.4]

arr2:
[[ 1.5  5.1  6.  12.  3.4]
 [ 1.5  5.1  6.  12.  3.4]
 [ 0.   1.   2.   3.   8. ]]

arr3:
[[ 1  5  6 12  3]
 [ 0  1  2  3  8]]

```

- There are a number of other functions for creating new arrays.

- As examples, `zeros` and `ones` create arrays of 0s or 1s, respectively, with a given length or shape.
- `empty` creates an array without initializing its values to any particular value.
- To create a higher dimensional array with these methods, pass a tuple for the shape:

Function	Description
<code>array</code>	Convert input data (list, tuple, array, or other sequence type) to an ndarray either by inferring a dtype or explicitly specifying a dtype; copies the input data by default
<code>asarray</code>	Convert input to ndarray, but do not copy if the input is already an ndarray
<code>arange</code>	Like the built-in <code>range</code> but returns an ndarray instead of a list
<code>ones</code> , <code>ones_like</code>	Produce an array of all 1s with the given shape and dtype; <code>ones_like</code> takes another array and produces a ones array of the same shape and dtype
<code>zeros</code> , <code>zeros_like</code>	Like <code>ones</code> and <code>ones_like</code> but producing arrays of 0s instead
<code>empty</code> , <code>empty_like</code>	Create new arrays by allocating new memory, but do not populate with any values like <code>ones</code> and <code>zeros</code>
<code>full</code> , <code>full_like</code>	Produce an array of the given shape and dtype with all values set to the indicated "fill value" <code>full_like</code> takes another array and produces a filled array of the same shape and dtype
<code>eye</code> , <code>identity</code>	Create a square $N \times N$ identity matrix (1s on the diagonal and 0s elsewhere)

```
arr1 = np.zeros((2,6))
print('\n arr1:\n', arr1)

arr2 = np.empty((2,3,2))
print('\n arr2:\n', arr2)

arr3 = np.arange(20)
print('\n arr3:\n', arr3)
```

```
arr1:
[[0. 0. 0. 0. 0. 0.]
 [0. 0. 0. 0. 0. 0.]]
```

```
arr2:
[[[0. 0.]
  [0. 0.]
  [0. 0.]]
```

```
[[0. 0.]
 [0. 0.]
 [0. 0.]]]
```

```
arr3:
[ 0  1  2  3  4  5  6  7  8  9 10 11 12 13 14 15 16 17 18 19]
```

Arithmetic with NumPy Arrays

- Arrays enable you to express batch operations on data without writing any forloops. NumPy users call this vectorization.
- Any arithmetic operations between equal-size arrays applies the operation element-wise.
- avoid using loops as much as you can while dealing with `ndarray`

Example 1: The following example show the difference between function based on a loop and another one based on ndarray `ufuncs`

```
import numpy as np

#define a standard function for calculating reciprocals using loop
def calc_rec(vect):
    res = np.empty(len(vect))
    for i in range(len(vect)):
        res[i] = 1.0 / vect[i]
    return res

vect = np.random.randint(1,20,5)
print("vect: \n", vect)
print("Reciprocals using our function: \n", calc_rec(vect))
print("Reciprocals using numpy ufuncs: \n", 1/vect)

vect:
[ 2  3  9  6 14]
Reciprocals using our function:
[0.5      0.33333333 0.11111111 0.16666667 0.07142857]
Reciprocals using numpy ufuncs:
[0.5      0.33333333 0.11111111 0.16666667 0.07142857]

# compare the approaches in terms of execution time

big_vect = np.random.randint(1,20,1000000)

%timeit calc_rec(big_vect)
%timeit (1/big_vect)

1 loop, best of 5: 2.22 s per loop
1000 loops, best of 5: 1.54 ms per loop
```

- Vectorized operations in NumPy are implemented via ufuncs.
- ufuncs's main purpose is to quickly execute repeated operations on values in NumPy arrays without needing loops.
- Ufuncs are extremely flexible (operations can be performed between scalar/ array and array/array).

Example 2:

```

import numpy as np

arr1 = np.ones((2,3))/(3 * np.random.randn(2,3)) #scaler/array and array/array operations
arr2 = np.power(np.reshape(arr1,(1,6), order='F'), 2)
arr3 = arr1+arr1
arr4 = arr3<0

print(' arr1:\n', arr1)
print('\n arr2:\n', arr2)
print('\n arr3:\n', arr3)
print('\n arr4:\n', arr4)

arr1:
[[-0.58760428 -4.14201903 -2.68074686]
 [-1.95918352  0.24392213 -0.3310625  ]]

arr2:
[[ 0.34527879  3.83840008 17.15632164  0.05949801  7.18640374
 0.10960238]]

arr3:
[[-1.17520856 -8.28403806 -5.36149373]
 [-3.91836705  0.48784426 -0.662125  ]]

arr4:
[[ True  True  True]
 [ True False  True]]

```

Example 3: Here are some other common ufuncs used with Numpy arrays. Each of these arithmetic operations are simply wrappers around specific functions built into NumPy; for example, the `+` operator is a wrapper for the `numpy.add` function

```

x = np.arange(4)
print("x      =", x)
print("x + 5   =", x + 5)
print("x - 5   =", x - 5)
print("x * 2   =", x * 2)
print("x / 2   =", x / 2)
print("x // 2  =", x // 2) # floor division
print("-x     =", -x)
print("abs(-x)=", abs(-x))
print("x ** 2  =", x ** 2)
print("x % 2   =", x % 2)
print("np.multiply.outer(x, x)", np.multiply.outer(x, x)) #vwT

x      = [0 1 2 3]
x + 5   = [5 6 7 8]
x - 5   = [-5 -4 -3 -2]
x * 2   = [0 2 4 6]

```

```

x / 2 = [0.  0.5 1.  1.5]
x // 2 = [0 0 1 1]
-x      = [ 0 -1 -2 -3]
abs(-x)= [0 1 2 3]
x ** 2 = [0 1 4 9]
x % 2   = [0 1 0 1]
np.multiply.outer(x, x) [[0 0 0 0]
 [0 1 2 3]
 [0 2 4 6]
 [0 3 6 9]]

# Trigonometric functions
theta = np.round(np.linspace(0, np.pi, 4), 2)
print("theta      = ", theta)
print("sin(theta) = ", np.sin(theta) )
print("cos(theta) = ", np.cos(theta) )
print("tan(theta) = ", np.tan(theta) )

theta      = [0.  1.05 2.09 3.14]
sin(theta) = [0.  0.86742323 0.86821458 0.00159265]
cos(theta) = [ 1.  0.49757105 -0.49618891 -0.99999873]
tan(theta) = [ 0.00000000e+00  1.74331531e+00 -1.74976619e+00 -
 1.59265494e-03]

#Exponents and logarithms
x = [1, 2, 3]
print("x          =", x)
print("e^x         =", np.exp(x))
print("2^x         =", np.exp2(x))
print("3^x         =", np.power(3, x))

print("ln(x)       =", np.log(x))
print("log2(x)      =", np.log2(x))
print("log10(x)     =", np.log10(x))

x          = [1, 2, 3]
e^x        = [ 2.71828183  7.3890561 20.08553692]
2^x        = [2.  4.  8.]
3^x        = [ 3  9 27]
ln(x)      = [0.  0.69314718 1.09861229]
log2(x)    = [0.  1.  1.5849625]
log10(x)   = [0.  0.30103  0.47712125]

```

Specialized ufuncs:

- NumPy has many more ufuncs available, look through the NumPy documentation to learn more.
- Another excellent source for more specialized ufuncs is the submodule `scipy.special`

```
from scipy import special
```

```

# Gamma functions (generalized factorials) and related functions
x = [1, 3, 5]
print("gamma(x)      =", special.gamma(x))
print("ln|gamma(x)| =", special.gammaln(x))
print("beta(x, 2)    =", special.beta(x, 2))

# Error function (integral of Gaussian), its complement, and its
inverse
x = np.array([0, 0.3, 0.7, 1.0])
print("\nerf(x)      =", special.erf(x))
print("erfc(x)     =", special.erfc(x))
print("erfinv(x)    =", special.erfinv(x))

gamma(x)      = [ 1.  2. 24.]
ln|gamma(x)|  = [0.          0.69314718 3.17805383]
beta(x, 2)    = [0.5          0.08333333 0.03333333]

erf(x)      = [0.          0.32862676 0.67780119 0.84270079]
erfc(x)     = [1.          0.67137324 0.32219881 0.15729921]
erfinv(x)    = [0.          0.27246271 0.73286908          inf]

```

Indexing and Slicing

- There are many ways you may want to select a subset of data or individual elements.
- One-dimensional ndarrays simply act similarly to Python lists.
- With Numpy, data is not copied, and any modifications to a view will be reflected in the source array. to copy a slice of an ndarray instead of a view, you will need `.copy`

```

import numpy as np

arr = np.arange(5) * 2
print('arr:      ', arr)
print('arr[2]:   ', arr[2])
print('arr[1:4]: ', arr[1:4])
arr[1:4] = 1
print('arr:      ', arr)

arr:      [0 2 4 6 8]
arr[2]:   4
arr[1:4]: [2 4 6]
arr:      [0 1 1 1 8]

# data modification
arr1 = np.arange(5) + 2
arr2 = arr1[1:4]
arr2[:] = 0
print('arr2:      ', arr2)
print('arr1:      ', arr1)
arr2 = arr1[1:4].copy() #this solves the problem of modifyng the
original data
arr2[:] = 1

```

```
print('-----')
print('arr2:      ', arr2)
print('arr1:      ', arr1)
```

```
arr2:      [0 0 0]
arr1:      [2 0 0 0 6]
-----
arr2:      [1 1 1]
arr1:      [2 0 0 0 6]
```

- With higher dimensional arrays, indices are no longer scalars but rather one-dimensional arrays.
- pass a comma-separated list of indices to select individual elements

```
arr2d = np.array([[1, 2, 3], [4, 5, 6], [7, 8, 9]])
print('arr2d[2]:      ', arr2d[2])
print('arr2d[3,2]:     ', arr2d[2,1])

arr3d = np.array([[[1, 2, 3], [4, 5, 6]], [[7, 8, 9], [10, 11, 12]]])

arr3d_2 = arr3d[1]
print('\narr3d_2:\n      ', arr3d_2)
print('\narr3d[1,1]:\n      ', arr3d[1,1])
print('\narr3d[1,1,2]:\n      ', arr3d[1,1,2])
```

```
arr2d[2]:      [7 8 9]
arr2d[3,2]:     8
```

```
arr3d_2:
[[ 7  8  9]
 [10 11 12]]
```


```
arr3d[1,1]:
[10 11 12]
```

```
arr3d[1,1,2]:
12
```

```
import numpy as np
import matplotlib.pyplot as plt
```

```
#Let's suppose that we have the following plotted profit of 4 days
profit = np.linspace(0,4,100)**2
profit = profit + np.random.randn(100)
plt.plot(np.linspace(0,4,100), profit)
```

```
[<matplotlib.lines.Line2D at 0x7f1e45ed2a10>]
```


```
#we want to plot the profit of day 2
fig = plt.figure(figsize=(10, 3))
profit_day2 = profit[25:50]

fig.add_subplot(1,2,1) # plot the slice
plt.plot(np.linspace(0,1,25), profit_day2)

fig.add_subplot(1,2,2) #plot the accumative sum
plt.plot(np.linspace(0,1,25), profit_day2.cumsum())

#show some statistics
print("max profit: ", profit_day2.max())
print("min profit: ", profit_day2.min())
print("mean profit: ", profit_day2.mean())
print("std profit: ", profit_day2.std())

max profit:  4.41976198733571
min profit:  -0.9933785966128232
mean profit:  2.117912064412047
std profit:  1.4166553214281168
```



Boolean Indexing Logical operators can be used to slice Numpy arrays *Example:* suppose that the following students recieved marks as follow

	Khaled	Selma	Djaber
PSD	10	12	9
DL	15	10	11
English	12	12	14

```
students = np.array(['Khaled', 'Selma', 'Djaber'])
modules = np.array(['PSD', 'DL', 'English'])
marks=np.array([[10,12,9], [15,10,11], [12,12,14]])

print("--students == 'Selma':", students == 'Selma')
print("--marks[modules == 'DL']:", marks[modules == 'DL'])
print("--marks[:,students == 'Selma']:", marks[:,students == 'Selma'])
print("--marks[:,students != 'Selma']:", marks[:,students != 'Selma'])
print("--marks[modules == 'PSD',students == 'Djaber']:", marks[modules == 'PSD',students == 'Djaber'])
marks[marks>10] = marks[marks>10] + 2

--students == 'Selma':      [False  True False]
--marks[modules == 'DL']:    [[15 10 11]]
--marks[:,students == 'Selma']: [[12]
[10]
[12]]
--marks[:,students != 'Selma']: [[10  9]
[15 11]
[12 14]]
--marks[modules == 'PSD',students == 'Djaber']: [9]
```

Fancy Indexing

- indexing using integer arrays.
- Fancy indexing, unlike slicing, always copies the data into a new array.

```

data = np.arange(20).reshape((5, 4))
print("---data:\n", data)
print("---data[[1,3,4]]:\n", data[[1,3,4]])
print("---data[[-4,-2,-1]]:\n", data[[-4,-2,-1]])
print("---data[[-8,-6,-1], [0,2,-1]]\n:", data[[-4,-2,-1], [0,2,-1]])

---data:
[[ 0  1  2  3]
 [ 4  5  6  7]
 [ 8  9 10 11]
 [12 13 14 15]
 [16 17 18 19]]
---data[[1,3,4]]:
[[ 4  5  6  7]
 [12 13 14 15]
 [16 17 18 19]]
---data[[-4,-2,-1]]:
[[ 4  5  6  7]
 [12 13 14 15]
 [16 17 18 19]]
---data[[-8,-6,-1], [0,2,-1]]
: [ 4 14 19]

```

Aggregations and sorting

Aggregation:

- When faced with a large amount of data, a first step is to compute summary statistics
- The most common summary statistics are the mean and standard deviation, which allow you to summarize the "typical" values in a dataset.
- Other aggregates are useful as well (the sum, product, median, minimum and maximum, quantiles, etc.)

```

import numpy as np
import matplotlib.pyplot as plt

#generate profit per day data
profit = 1000*np.random.randn(20)**2
plt.plot(np.arange(1,21, 1), profit)
plt.xticks(np.arange(1,21, 1))

#some aggeragations
mean = np.mean(profit);
std = np.std(profit);
median = np.median(profit)
print("mean: {mean:.2f}, standaed deviation: {std:.2f}, and med:
{mean:.2f}".format(mean=mean, std=std, med = median))


#other aggregations
min = np.min(profit)

```

```
x_min = np.argmin(profit) + 1
plt.plot(x_min, min, 'r*')
```

```
max = np.max(profit)
x_max = np.argmax(profit) + 1
plt.plot(x_max, max, 'g*');
```

mean: 994.76, standaed deviation: 1323.06, and med: 994.76



For multi dimentional arrays, the aggregtions are estimated over a specefied axis

```
arr = np.array([[0, 1, 2], [3, 4, 5], [6, 7, 8]])
print('arr:\n', arr)
arr2 = arr.cumsum(axis=0)
print('arr2:\n', arr2)
np.max(arr, 0)
```

```
arr:
[[0 1 2]
 [3 4 5]
 [6 7 8]]
arr2:
[[ 0  1  2]
 [ 3  5  7]
 [ 9 12 15]]
```

```
array([6, 7, 8])
```

Sorting

- This subsection covers algorithms related to sorting values in NumPy arrays.
- These algorithms are of fundamental topics in introductory computer science courses (insertion sorts, selection sorts, merge sorts, quick sorts, bubble sorts, etc.)
- NumPy arrays can be sorted in-place with the `sort` method of order $O(N \log N)$.
- `argsort` is a related function, which instead returns the indices of the sorted elements.

```
import numpy as np
#sorting one dimension array
arr = np.abs(np.random.randn(5))
sorted_arr = np.sort(arr)
sorted_indices = np.argsort(arr)

print('arr: ', arr)
print('sorted array: ', sorted_arr)
print('sorted indices:', sorted_indices)
print('\n\n')

#sorting n dimension array
arr = np.random.randint(0, 10, (3, 4))
sorted_arr_over_rows = np.sort(arr, axis=0)
sorted_arr_over_cols = np.sort(arr, axis=1)
print('arr:\n', arr)
print('sorted array over rows:\n', sorted_arr_over_rows)
print('sorted array over columns:\n', sorted_arr_over_cols)
print('\n\n')
```

arr: [0.84778948 0.06092322 0.62543832 1.57583266
0.59100793]
sorted array: [0.06092322 0.59100793 0.62543832 0.84778948
1.57583266]
sorted indices: [1 4 2 0 3]


```
arr:
[[6 3 4 1]
 [0 6 7 7]
 [9 6 0 7]]
sorted array over rows:
[[0 3 0 1]
 [6 6 4 7]
 [9 6 7 7]]
sorted array over columns:
[[1 3 4 6]
 [0 6 7 7]
 [0 6 7 9]]
```

- sometimes we are interested in finding and sorting the k smallest values in the array.

- `np.partition` takes an array and a number K , and produces a new array with the smallest K values to the left of the partition, and the remaining values to the right, in arbitrary order.

```
x = np.array([7, 2, 3, 1, 5, 8, 4])
np.partition(x, 3)

array([2, 1, 3, 4, 5, 8, 7])
```

Example: k-Nearest Neighbors:

- The k-nearest neighbors (KNN) is a simple supervised machine learning algorithm that can be used to solve classification and regression problems.
- KNN algorithm assumes that similar things exist in close proximity. In other words, similar things are near to each other.
- For a given point in the space (image, packet, word, etc), KNN aims at finding the most similar points(i.o.w., finding neighbors after sorting).

```
import numpy as np
import matplotlib.pyplot as plt
import warnings
warnings.filterwarnings('ignore')
fig = plt.figure(figsize=(4, 4))

#Generate Data
Birds = np.random.multivariate_normal([12, 1], [[1,2],[2,1]], 10).T
Reptiles = np.random.multivariate_normal([1, 12], [[1,2],[2,1]], 10).T
Insects = np.random.multivariate_normal([1, 1], [[1,2],[2,1]], 10).T

print("Birds:\n", Birds)
print("Repriles:\n", Reptiles)
print("Insects:\n", Insects)

Birds:
[[12.90167044 14.07200497 10.54703882 13.19117433 13.0424816
13.02076662
10.96055607 12.01516425 12.16171476 11.99964493]
[ 1.25530601  2.04573506  0.60441732  2.6614238  3.0181768
0.56147608
0.81673676  0.07888133 -0.38632499  0.20938731]]
Repriles:
[[ 0.71545065  3.765862  1.30666699  2.62239463  2.33971954 -
1.54288147
2.25728551  3.49278675  4.15338558  1.99934912]
[10.2127597 14.84638237 11.23207262 15.5107982 13.85921634
11.86678895
13.5720774 13.1557101 13.34748752 11.82452219]]
Insects:
[[ 1.45795908  0.35836242  1.83584422  0.40669918  2.70512557
2.02988764
-0.85182524 -0.31542024  1.63233144 -3.49929197]
```


```
[ -0.25284019  1.10457218  0.42908635  3.87453712  0.80574787
 1.98689088
 -0.675192   -1.41049149 -0.65751082 -0.8707069  ]]
```

<Figure size 288x288 with 0 Axes>

To get an idea of how these points look, let's quickly scatter plot them.

```
plt.scatter(Birds[0], Birds[1], marker='o', label='Birds' )
plt.scatter(Reptiles[0], Reptiles[1], marker='x' , label='Reptiles')
plt.scatter(Insects[0], Insects[1], marker='v' , label='Insects');
plt.legend()
```

<matplotlib.legend.Legend at 0x7f1e46ca9250>



Let's suppose that we have some new data we want to recognize

```
new_point = np.random.randint(-2,14,2)
```

```
#plot clases
```


```
plt.scatter(Birds[0], Birds[1], marker='o', label='Birds' )
plt.scatter(Reptiles[0], Reptiles[1], marker='x' , label='Reptiles')
plt.scatter(Insects[0], Insects[1], marker='v' , label='Insects');
```

```
#plot the new point
```

```
plt.plot(new_point[0], new_point[1],   "**r", markersize=12, label='New
data')
```

```
plt.legend()
```

```
<matplotlib.legend.Legend at 0x7f1e466fb050>
```



The first step is to calculate the distance between this point all point of the three classes.

```
#merge the data into one single array
```

```
K = 3
```

```
data = np.concatenate( (Birds,Reptiles,Insects), axis=1)
```

```
distances = np.sum((data - new_point[:,np.newaxis])**2, axis=0)
```

```
sorted_points = np.argsort(distances, K+1)
```

```
print(sorted_points)
```

```
k_nearest = sorted_points[0:K]
```

```
#plot the result
```

```
plt.scatter(Birds[0], Birds[1], marker='o', label = 'Birds' )
```

```
plt.scatter(Reptiles[0], Reptiles[1], marker='x', label = 'Reptiles' )
```

```
plt.scatter(Insects[0], Insects[1], marker='v' , label = 'Insects');
```

```
#plot the point
```

```
plt.plot(new_point[0], new_point[1],  "**r", markersize=12, label  
='New Data')
```

```
plt.legend()
```

```
#plot lines to the closest points
```


```
for i in range(3):
```

```
    plt.plot(*zip(new_point, data[:,k_nearest[i]]), c='k')
```



```
#print the predicted class
unique, counts = np.unique(np.uint8(k_nearest/10), return_counts=True)
c = unique[np.argmax(counts)]
print("The class of the new data is: \x1b[1;31m", (['Birds',
'Reptiles', 'Insects'])[c], "\x1b[0m")

[18 17 11 14 16 19 13 12 10 15 23  3 24  6  1 25  4  2  0  9 20 21 22
 8
 7  5 26 27 28 29]
The class of the new data is: Reptiles
array([1], dtype=uint8)
```



1.6: Data manipulation with Pandas

- Pandas is a package built upon NumPy providing an efficient implementation of a **DataFrame**.
- DataFrames are essentially n-dimensional arrays with row/column labels, and often with heterogeneous types and/or missing data.
- While numpy serves its purpose very well, its limitations become clear when we need more flexibility (e.g., attaching labels to data, working with missing data, etc.) and when attempting operations that do not map well to element-wise broadcasting (e.g., groupings, pivots, etc.).

```
import pandas as pd
```

Pandas Series

To create a data **Series**

```
data = pd.Series([0.25, 0.5, 0.75, 1.0])
print("values : ", data.values)
print("indices: ", data.index)

values : [0.25 0.5 0.75 1. ]
indices: RangeIndex(start=0, stop=4, step=1)
```

data can be accessed as with numpy

```
print("data[1]:      ", data[1])
print("data[1:-1]: ", data[1:-1].values)

data[1]:      0.5
data[1:-1]: [0.5 0.75]
```

Numpy Array has an implicitly defined integer index used to access the values, the Pandas Series has an explicitly defined index associated with the values.

```
data = pd.Series([0.25, 0.5, 0.75, 1.0],
                  index=['a', 'b', 'c', 'd'])
print("data[1]:      ", data['a'])
print("data[-1:-3]: ", data['b':'d'].values)

data[1]:      0.25
data[-1:-3]: [0.5 0.75 1. ]
```

- Pandas Series are a bit like Python dictionary.
- A dictionary is a structure that maps **arbitrary** keys to a set of arbitrary values,
- Series is a structure which maps **typed** keys to a set of typed values.
- This typing makes Series more efficient than a Python lists for certain operations.

```
pop_dict = {'Constantine': 938475,
            'Biskra': 547137,
            'Ouargla': 311337,
            'Eloued': 647548,
            'Adrar': 261258}

population = pd.Series(pop_dict)
print('population[Ouargla] : ',
      population['Ouargla'])
print("population[['Biskra', 'Eloued', 'Adrar']] : ",
      population[['Biskra', 'Eloued', 'Adrar']].values)
print('population[0:2] : ',
      population[0:2].values)
print('population[(population>500000) & (population<700000)] : ',
      population[(population>500000) & (population<700000)].index.values)

population[Ouargla] : 311337
population[['Biskra', 'Eloued', 'Adrar']] : [547137 647548
261258]
```

```
population[0:2] : [938475
547137]
population[(population>500000) & (population<700000)] : ['Biskra'
'Eloued']
```

Constructing Series: The main instruction for constructing series is `pd.Series(data, index=index)`. Other variantes of this isntruction can be used:

```
pd.Series([5.6, 13.1, 19])

0    2
1    4
2    6
dtype: int64

pd.Series('Hello', index=[10, 20, 50])

10    Hello
20    Hello
50    Hello
dtype: object

pd.Series({2:'X', 5:'Y', 1:'Z'})

2    X
5    Y
1    Z
dtype: object

pd.Series({2:'a', 1:'b', 3:'c', 9:'d'}, index=[3, 2, 9])

3    c
2    a
9    d
dtype: object
```

Indexers: loc, iloc, and ix: if your Series has an explicit integer index, the *indexing* uses explicit while *slicing* use the implicit indexing.

```
data = pd.Series(['a', 'b', 'c', 'd'], index=[2, 3, 5, 6])
print('data[1] :', data[2]) # explicit index when indexing
print('data[1:3] :', data[2:4].values) # implicit index when slicing

data[1] : a
data[1:3] : ['c' 'd']
```

`loc` attribute allows indexing and slicing that always references the explicit index:

```
print('data.loc[1] :', data.loc[2])
print('data.loc[1:3] :', data.loc[2:5].values)
```

```
data.loc[1]    : a
data.loc[1:3] : ['a' 'b' 'c']
```

iloc attribute allows implicit indexing and slicing

```
print('data.iloc[1]    : ', data.iloc[1])
print('data.iloc[1:3] : ', data.iloc[1:3].values)

data.iloc[1]    : b
data.iloc[1:3] : ['b' 'c']
```

- One guiding principle of Python code is that "explicit is better than implicit."
- Use both to make code easier to read and understand, and to prevent subtle bugs due to the mixed indexing/slicing convention.

Pandas DataFrame

- `DataFrame` is an analog of a two-dimensional array with both flexible row indices and flexible column names.
- You can think of a `DataFrame` as a sequence of aligned `Series` objects.

Example:

```
import pandas as pd

pop_dict    = {'Constantine': 938475, 'Biskra': 547137, 'Ouargla': 311337, 'Eloued': 647548, 'Adrar': 261258}
area_dict   = {'Constantine': 2187, 'Biskra': 9576, 'Ouargla': 194552, 'Eloued': 45738, 'Adrar': 254471}
population  = pd.Series(pop_dict)
area        = pd.Series(area_dict)

#construct a Dataframe
wilayas = pd.DataFrame({'population': population, 'area': area})
print("indices: ", wilayas.index)
print("columns: ", wilayas.columns)
wilayas

indices:  Index(['Constantine', 'Biskra', 'Ouargla', 'Eloued', 'Adrar'], dtype='object')
columns:  Index(['population', 'area'], dtype='object')
```

	population	area
Constantine	938475	2187
Biskra	547137	9576
Ouargla	311337	194552
Eloued	647548	45738
Adrar	261258	254471

we can pick one column from the `DataFrame`

```
wilayas['area']
```

Constantine	2187
Biskra	9576
Ouargla	194552
Eloued	45738
Adrar	254471

Name: area, dtype: int64

In case some keys are missing, **Pandas** will fill them in with NaN

```
pd.DataFrame([{'a': 1, 'b': 2, 'c': 5}, {'b': 3, 'c': 4, 'd': 12}])
```

	a	b	c	d
0	1.0	2	5	NaN
1	NaN	3	4	12.0

We can create `DataFrame`, from any numpy array, with any specified column and index names

```
import numpy as np

pd.DataFrame(np.random.rand(5, 3),
             columns=['income', 'outcome', 'lose'],
             index=['day 1', 'day 2', 'day 3', 'day 4', 'day 5'])
```

	income	outcome	lose
day 1	0.175215	0.437155	0.053180
day 2	0.632874	0.782488	0.187896
day 3	0.393341	0.215869	0.123302
day 4	0.163537	0.711649	0.676054
day 5	0.695072	0.392600	0.372434

Data Indexing and Slicing

```
import pandas as pd

pop_dict = {'Constantine': 938475, 'Biskra': 547137, 'Ouargla': 311337, 'Eloued': 647548, 'Adrar': 261258}
area_dict = {'Constantine': 2187, 'Biskra': 9576, 'Ouargla': 194552, 'Eloued': 45738, 'Adrar': 254471}

wilayas = pd.DataFrame({'population': pop_dict, 'area': area_dict})
#add a column
wilayas['density'] = wilayas['population']/ wilayas['area']
wilayas
```

	population	area	density
Constantine	938475	2187	429.115226
Biskra	547137	9576	57.136278
Ouargla	311337	194552	1.600277

Eloued	647548	45738	14.157768
Adrar	261258	254471	1.026671

One can transpose the full DataFrame to swap rows and columns

wilayas.T

	Constantine	Biskra	...	Eloued
Adrar				
population	938475.000000	547137.000000	...	647548.000000
261258.000000				
area	2187.000000	9576.000000	...	45738.000000
254471.000000				
density	429.115226	57.136278	...	14.157768
1.026671				

[3 rows x 5 columns]

Indexing & slicing

```
print('wilayas.values[0]      :', wilayas.values[0])
print("\nwilayas['density']    :\n", wilayas['density'])
print("\nwilayas.iloc[0:3, 1:]  :\n", wilayas.iloc[0:3, 1:]) #
slicing the array as if it is a simple NumPy array using loc
print("\nwilayas.loc['Ouargla':, : 'area'] :\n",
wilayas.loc['Ouargla':, : 'area']) # slicing using loc
print("\nwilayas.loc['Ouargla':, : 'area'] :\n",
      wilayas.loc[wilayas.density > 50, ['population', 'area']])
#fancy indexing
```

wilayas.values[0] : [9.38475000e+05 2.18700000e+03 4.29115226e+02]

wilayas['density'] :
Constantine 429.115226
Biskra 57.136278
Ouargla 1.600277
Eloued 14.157768
Adrar 1.026671
Name: density, dtype: float64

wilayas.iloc[0:3, 1:] :
 area density
Constantine 2187 429.115226
Biskra 9576 57.136278
Ouargla 194552 1.600277

wilayas.loc['Ouargla':, : 'area'] :
 population area
Ouargla 311337 194552

Eloued	647548	45738
Adrar	261258	254471

```
wilayas.loc['Ouargla':, : 'area'] :
```

	population	area
Constantine	938475	2187
Biskra	547137	9576

- Any of these indexing/slicing approaches may also be used to set or modify values;
- The `ix` indexer allows a hybrid of these two approaches. Try it yourselves.

Operations

- Pandas inherits much functionalities from NumPy

```
import numpy as np

ds = pd.Series({'a':10, 'b':2, 'c':5, 'd':7})
df = pd.DataFrame({'x':ds, 'y':(ds + ds)**2/3})
np.exp(data_1)
np.sin(df * np.pi / 4)

32.66666666666667
```

Operating on Null Values:

```
import numpy as np
data = pd.Series([1, np.nan, 'hello', None])
print(data.isnull().values)
print(data[data.notnull()].values)
data.dropna()

[False  True  False  True]
[1 'hello']

0      1
2    hello
dtype: object
```

Hierarchical Indexing

- Often it is useful to go beyond this and store higher-dimensional data with more than one or two keys.
- a common pattern in practice is to make use of hierarchical indexing to incorporate multiple index levels within a single index.

```
import pandas as pd

pop_dict = {('Constantine', 2010): 538475, ('Constantine', 2020):
938475,
            ('Biskra', 2010): 247137, ('Biskra', 2020): 247137,
```

```

                ('Ouargla', 2010): 111337, ('Ouargla', 2020): 311337}
# area_dict = {'Constantine': 2187, 'Biskra': 9576, 'Ouargla':
194552 }

```

```

wilayas = pd.Series(pop_dict)
wilayas.index.names=('wilaya', 'year')#Sometimes it is convenient to
name the levels of the MultiIndex
print("wilayas.loc[['Biskra', 2010],('Constantine', 2020)]: \n",
wilayas.loc[['Biskra', 2010],('Constantine', 2020)], '\n')
wilayas

```

```

wilayas.loc[['Biskra', 2010],('Constantine', 2020)]:
wilaya      year
Biskra      2010    247137
Constantine 2020    938475
dtype: int64

```

```

wilaya      year
Constantine 2010    538475
            2020    938475
Biskra      2010    247137
            2020    247137
Ouargla     2010    111337
            2020    311337
dtype: int64

```

How to select all populations in 2010?

```

print(wilayas.loc[[i for i in wilayas.index if i[1] == 2010]],'\n')
print(wilayas[:,2010],'\n')

```

```

Constantine 2010    538475
Biskra      2010    247137
Ouargla     2010    111337
dtype: int64

```

```

Constantine    538475
Biskra         247137
Ouargla        111337
dtype: int64

```

The `unstack()` method will quickly convert a multiply indexed `Series` into a conventionally indexed `DataFrame`. Naturally, the `stack()` method provides the opposite operation

```

wilayas.unstack()

          2010    2020
Biskra    247137  247137

```


Constantine	538475	938475
Ouargla	111337	311337

why would we would bother with hierarchical indexing at all?

- to use multi-indexing to represent two-dimensional data within a one-dimensional Series,
- we can also use it to represent data of three or more dimensions in a Series or DataFrame.
- Each extra level in a multi-index represents an extra dimension of data;

```
df = pd.DataFrame({'total': wilayas, 'under 10': wilayas.values//5})
df
```

		total	under 10
Constantine	2010	538475	107695
	2020	938475	187695
Biskra	2010	247137	49427
	2020	247137	49427
Ouargla	2010	111337	22267
	2020	311337	62267

```
(df['under 10']/df['total']).unstack()
```

	2010	2020
Biskra	0.199998	0.199998
Constantine	0.200000	0.200000
Ouargla	0.199996	0.199999

Combining Datasets: Concat and Append

Some of the most interesting studies of data come from combining different data sources

```
import numpy as np
import pandas as pd

#with numpy
x = [[1, 2], [3, 4]]
print('np.concatenate([x, x], axis=1):\n', np.concatenate([x, x],
axis=1))

#with pandas Series
ser1 = pd.Series(['A', 'B'], index=[1, 2])
ser2 = pd.Series(['D', 'E'], index=[4, 5])
print('\nnpd.concat([ser1, ser2]):\n', pd.concat([ser1, ser2]))

#with pandas Dataframes
df1 = pd.DataFrame(({1,2}, {3,4}), index=[0,1], columns=('C1', 'C2'))
df2 = pd.DataFrame(({11,22}, {33,44}), index=[0,1], columns=('C1',
```

```

'C2'))
print('\nnpd.concat([df1, df2]) :\n', pd.concat([df1, df2]))#concat
over rows
print('\nnpd.concat([df1, df2], axis=1) :\n', pd.concat([df1, df2],
axis=1)) #concat over columns
print('\nnpd.concat([df1, df2]) :\n', pd.concat([df1, df2],
keys=('df1', 'df2'))) #specify multindex for Hierarchical Indexing

np.concatenate([x, x], axis=1):
[[1 2 1 2]
 [3 4 3 4]]

pd.concat([ser1, ser2]):
1      A
2      B
4      D
5      E
dtype: object

pd.concat([df1, df2]) :
   C1  C2
0    1   2
1    3   4
0   11  22
1   33  44

pd.concat([df1, df2], axis=1) :
   C1  C2  C1  C2
0    1   2  11  22
1    3   4  33  44

pd.concat([df1, df2]) :
   C1  C2
df1 0    1   2
    1    3   4
df2 0   11  22
    1   33  44

```

Concatenation with joins

- we were mainly concatenating DataFrames with shared column names. *In practice, data from different sources might have different sets of column names
- `pd.concat` offers several options in this case.

```

df1 = pd.DataFrame(({1,2}), (3,4}), index=[0,1], columns=('C1', 'C2'))
df2 = pd.DataFrame(({1,2}), (3,4}), index=[0,1], columns=('C1', 'C3'))

print('pd.concat([df1, df2]):\n', pd.concat([df1, df2])) #inner
concatenation

```

```
print("\nnpd.concat([df1, df2], join='inner'):\n", pd.concat([df1, df2], join='inner')) #inner concatenation (union)
```

```
pd.concat([df1, df2]):
```

	C1	C2	C3
0	1	2.0	NaN
1	3	4.0	NaN
0	1	NaN	2.0
1	3	NaN	4.0

```
pd.concat([df1, df2], join='inner'):
```

	C1
0	1
1	3
0	1
1	3

GroupBy: Split, Apply, Combine

- Simple aggregations can give you a flavor of your dataset
- Often we would prefer to aggregate conditionally on some label or index
- This is implemented in the so-called **groupby** operation.

```
df = pd.DataFrame({'key': ['A', 'B', 'C', 'A', 'B', 'C'],
                   'data': range(6)}, columns=['key', 'data'])
```

```
df
```

	key	data
0	A	0
1	B	1
2	C	2
3	A	3
4	B	4
5	C	5

```
df.groupby('key').mean()
```

```
df.groupby('key')['data'].mean()
```

```
key
```

A	1.5
B	2.5
C	3.5

```
Name: data, dtype: float64
```

GroupBy object supports direct iteration over the groups

```
import pandas
import seaborn as sns
```

```
penguins = sns.load_dataset('penguins')
penguins.head()
```

	species	island	bill_length_mm	bill_depth_mm	flipper_length_mm
0	Adelie	Torgersen	39.1	18.7	181.0
1	Adelie	Torgersen	39.5	17.4	186.0
2	Adelie	Torgersen	40.3	18.0	195.0
3	Adelie	Torgersen	NaN	NaN	NaN
4	Adelie	Torgersen	36.7	19.3	193.0

	body_mass_g	sex
0	3750.0	Male
1	3800.0	Female
2	3250.0	Female
3	NaN	NaN
4	3450.0	Female

```
penguins.groupby('island').mean()
```

	bill_length_mm	bill_depth_mm	flipper_length_mm
Biscoe	45.257485	15.874850	209.706587
Dream	44.167742	18.344355	193.072581
Torgersen	38.950980	18.429412	191.196078

```
for (method, group) in penguins.groupby('island'):
    print("{0:30s} shape={1}".format(method, group.shape))
```

Biscoe	shape=(168, 7)
Dream	shape=(124, 7)
Torgersen	shape=(52, 7)

```
penguins.groupby('island')['body_mass_g'].describe()
```

	count	mean	std	...	50%	75%
max						
island				...		
Biscoe	167.0	4716.017964	782.855743	...	4775.0	5325.00
Dream	124.0	3712.903226	416.644112	...	3687.5	3956.25

```
4800.0
Torgersen    51.0  3706.372549  445.107940  ...  3700.0  4000.00
4700.0
```

```
[3 rows x 8 columns]
```

`aggregate()` method allows for even more flexibility. It can take a string, a function, or a list thereof, and compute all the aggregates at once.

```
penguins.groupby('island')['body_mass_g'].aggregate(['count', min,
np.median, max, np.var])
```

	count	min	median	max	var
island					
Biscoe	167	2850.0	4775.0	6300.0	612863.114133
Dream	124	2700.0	3687.5	4800.0	173592.315762
Torgersen	51	2900.0	3700.0	4700.0	198121.078431

```
penguins.groupby('island').aggregate({'max_body_mass_g': max,
'min_bill_depth_mm':min })
```

	body_mass_g	bill_depth_mm
island		
Biscoe	6300.0	13.1
Dream	4800.0	15.5
Torgersen	4700.0	15.9

transformation can return some transformed version of the full data to recombine.

```
penguins.groupby('island').transform(lambda x: x/ x.mean()).head()
```

```
/usr/local/lib/python3.7/dist-packages/ipykernel_launcher.py:1:
FutureWarning: Dropping invalid columns in DataFrameGroupBy.transform
is deprecated. In a future version, a TypeError will be raised. Before
calling .transform, select only columns which should be valid for the
transforming function.
```

```
"""Entry point for launching an IPython kernel.
```

	bill_length_mm	bill_depth_mm	flipper_length_mm	body_mass_g
0	1.003826	1.014682	0.946672	1.011771
1	1.014095	0.944143	0.972823	1.025261
2	1.034634	0.976700	1.019895	0.876868
3	NaN	NaN	NaN	NaN
4	0.942210	1.047239	1.009435	0.930829

The `apply()` method lets you apply an arbitrary function to the group results.

```
def norm_bodymass_by_sum_billdepth(x):
    # x is a DataFrame of group values
```

```

x['body_mass_g'] /= x['bill_depth_mm'].sum()**2
return x
penguins.groupby('island').apply(norm_bodymass_by_sum_billdepth).head(
)

```

	species	island	bill_length_mm	bill_depth_mm	flipper_length_mm
0	Adelie	Torgersen	39.1	18.7	181.0
1	Adelie	Torgersen	39.5	17.4	186.0
2	Adelie	Torgersen	40.3	18.0	195.0
3	Adelie	Torgersen	NaN	NaN	NaN
4	Adelie	Torgersen	36.7	19.3	193.0

	body_mass_g	sex
0	0.004245	Male
1	0.004302	Female
2	0.003679	Female
3	NaN	NaN
4	0.003905	Female

we can perform multidimensional grouping

```

penguins.groupby(['sex', 'island'])['body_mass_g'].mean().unstack()

```

island	Biscoe	Dream	Torgersen
sex			
Female	4319.375000	3446.311475	3395.833333
Male	5104.518072	3987.096774	4034.782609

The same operation can be executed using `pivot_table`

```

penguins.pivot_table('body_mass_g', 'sex', 'island')

```

island	Biscoe	Dream	Torgersen
sex			
Female	4319.375000	3446.311475	3395.833333
Male	5104.518072	3987.096774	4034.782609

we can build even more indices

```

penguins.pivot_table('body_mass_g', ['sex', 'species'], 'island')

```

island		Biscoe	Dream	Torgersen
sex	species			
Female	Adelie	3369.318182	3344.444444	3395.833333

	Chinstrap	NaN	3527.205882	NaN
	Gentoo	4679.741379	NaN	NaN
Male	Adelie	4050.000000	4045.535714	4034.782609
	Chinstrap	NaN	3938.970588	NaN
	Gentoo	5484.836066	NaN	NaN

even more column indices

```
import pandas as pd

mass = pd.qcut(penguins['body_mass_g'], 3)
penguins.pivot_table('body_mass_g', ['sex', 'species'], [mass, 'island'])
```

body_mass_g		(2699.999, 3700.0]	...	(4550.0, 6300.0]	
island		Biscoe	Dream	...	Dream
Torgersen					
sex	species			...	
Female	Adelie	3207.8125	3344.444444	...	NaN
NaN					
	Chinstrap	NaN	3446.428571	...	NaN
NaN					
	Gentoo	NaN	NaN	...	NaN
NaN					
Male	Adelie	3600.0000	3525.000000	...	4625.0
4687.5					
	Chinstrap	NaN	3487.500000	...	4800.0
NaN					
	Gentoo	NaN	NaN	...	NaN
NaN					

[6 rows x 9 columns]

`mean` is the default function for pivoting. However, aggregation method can be specified or even personalized.

```
penguins.pivot_table('body_mass_g', ['sex', 'species'], 'island',
aggfunc={sum})
```

		sum		
island		Biscoe	Dream	Torgersen
sex	species			
Female	Adelie	74125.0	90300.0	81500.0
	Chinstrap	NaN	119925.0	NaN
	Gentoo	271425.0	NaN	NaN
Male	Adelie	89100.0	113275.0	92800.0

Chinstrap	NaN	133925.0	NaN
Gentoo	334575.0	NaN	NaN

Example: Let's analyse the following data

```
!wget https://raw.githubusercontent.com/jakevdp/data-
CDCbirths/master/births.csv
--2022-11-02 11:34:39--
https://raw.githubusercontent.com/jakevdp/data-CDCbirths/master/births
.csv
Resolving raw.githubusercontent.com (raw.githubusercontent.com)...
185.199.108.133, 185.199.110.133, 185.199.111.133, ...
Connecting to raw.githubusercontent.com (raw.githubusercontent.com)|
185.199.108.133|:443... connected.
HTTP request sent, awaiting response... 200 OK
Length: 264648 (258K) [text/plain]
Saving to: 'births.csv'

births.csv      100%[=====>] 258.45K  --.-KB/s   in
0.03s

2022-11-02 11:34:39 (7.65 MB/s) - 'births.csv' saved [264648/264648]
```

```
import pandas as pd

data = pd.read_csv('births.csv')
#fill missings
data.fillna(method='ffill')
print("original data shape:", data.shape)

data.head()


original data shape: (15547, 5)
```

	year	month	day	gender	births
0	1969	1	1.0	F	4046
1	1969	1	1.0	M	4440
2	1969	1	2.0	F	4454
3	1969	1	2.0	M	4548
4	1969	1	3.0	F	4548

aggagte data based on number of births/year

```
sums = data.groupby('year')['births'].sum().plot()
sums

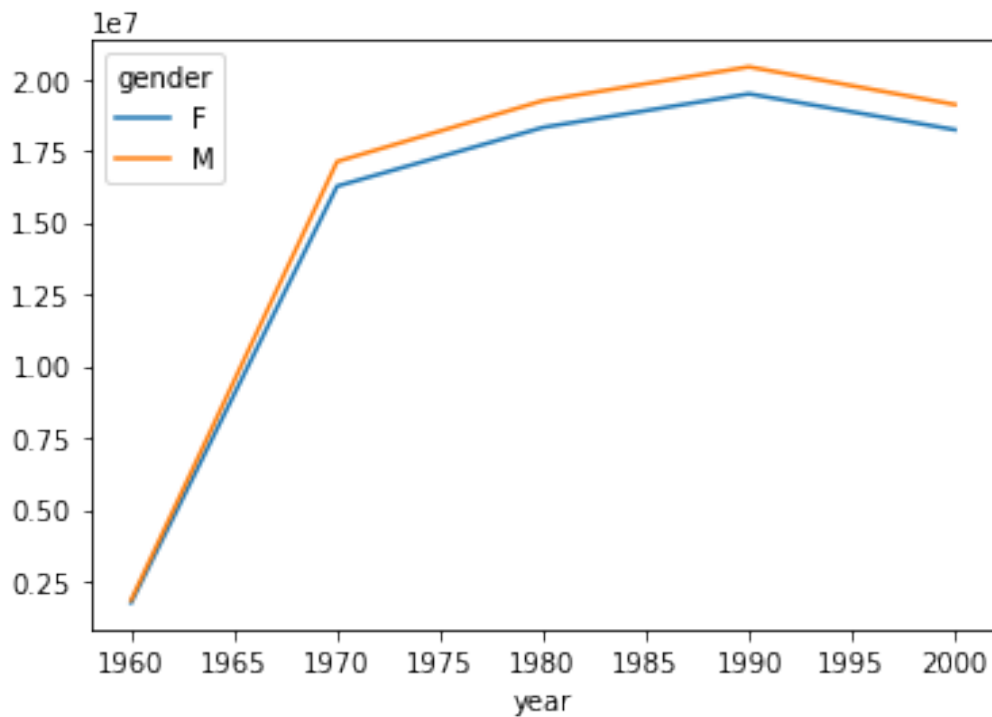
<matplotlib.axes._subplots.AxesSubplot at 0x7f9b85a11790>
```

what are the averages of births by gender in 5 equal periods?

```
decades = 10*(data['year']//10)
data.pivot_table('births', index=decades, columns='gender',
aggfunc='sum').plot()
data.pivot_table('births', 'gender', decades, aggfunc='sum')
```

year	1960	1970	1980	1990	2000
gender					
F	1753634	16263075	18310351	19479454	18229309
M	1846572	17121550	19243452	20420553	19106428



Let's see the births quantity per week days

```
import pandas as pd
data['dayofweek'] = pd.to_datetime(10000 * data.year +
                                   100 * data.month +
                                   data.day, format='%Y%m%d',
errors='coerce').dt.day_name()

data.groupby('dayofweek')['births'].sum().plot
<matplotlib.axes._subplots.AxesSubplot at 0x7f3feab86710>
```

